



ITCS413: Database Design
Project 5: Physical Database Design and Tuning

Airline: Star Airline

(May 2025)

Prepared by

6588070 Nakarin Phoorahong

6588096 Panipak Sittiprasert

6588183 Achiraya Mankham

Presented To

Asst.Prof. Dr. Charnyote Pluempitiwiriyawej

Faculty of Information and Communication Technology
Mahidol University

Table of contents

Database Application Requirements	1
Description of Required Transactions/Queries	1
Highlighted Selected Transactions/Queries	2
Final Conceptual Database Model	3
ER Diagram	3
Highlighted Portion Related to Selected Transactions/Queries	3
Final Logical Database Model	5
Relational Database Schema	5
Highlighted Portion Related to Selected Transactions/Queries	6
SQL Commands and Table Specifications	7
SQL Commands for Selected Transactions/Queries	7
Table Specifications (Number of Records, Record Sizes)	8
Pre-Improvement Transaction Analysis	9
Performance Results from RDBMS Query Analyzer	9
Post-Improvement Transaction Analysis	18
Performance Results after Optimization	20
Comparison with Pre-Improvement Results	30
Discussion	31
References	32

Database Application Requirements

Description of Required Transactions/Queries

Transaction 1: Count Total Passengers on a Specific Flight

- This transaction retrieves the number of passengers booked on a specific flight. It is vital for monitoring passenger load, optimizing flight resources, and ensuring efficient seat management.
- **Entities Involved:**
 - passenger
 - booking
 - passenger_booking
 - booking_flight
 - Flight

Transaction 2: List All Passengers on a Specific Flight

- This transaction retrieves the list of all passengers booked on a specific flight. It supports boarding management, security checklists, and passenger services.
- **Entities Involved:**
 - passenger
 - booking
 - passenger_booking
 - booking_flight
 - Flight

Transaction 3: Calculate Total Revenue from Completed Payments

- This transaction calculates the total revenue from all successfully completed passenger payments. It is critical for financial reporting and revenue management.
- **Entities Involved:** payment

Highlighted Selected Transactions/Queries

Transaction 1:

- **Highlighted Portions:**
 - Entities: Passenger, Booking, Flight
- **Relationships:**
 - Passenger ↔ Booking via passenger_booking (M:N)
 - Booking ↔ Flight via booking_flight (M:N)

Transaction 2:


- **Highlighted Portions:**
 - Entities: Passenger, Booking, Flight
- **Relationships:**
 - Passenger ↔ Booking via passenger_booking (M:N)
 - Booking ↔ Flight via booking_flight (M:N)

Transaction 3:

- **Highlighted Portions:**
 - Entities: Payment
- **Relationships:** None (single table used)


Final Conceptual Database Model

ER Diagram




Highlighted Portion Related to Selected Transactions/Queries


Transaction 1:



Transaction 2:




Transaction 3:



Final Logical Database Model

Relational Database Schema



Highlighted Portion Related to Selected Transactions/Queries

Transaction 1:

- passenger(passengerID, fullName, phoneNum, email)
- booking(bookingID)
- flight(flightID, depAirport, arrAirport, stops)
- passenger_booking(passengerID, bookingID)
- booking_flight(bookingID, flightID)

Transaction 2:

- passenger(passengerID, fullName, phoneNum, email)
- booking(bookingID)
- flight(flightID, depAirport, arrAirport, stops)
- passenger_booking(passengerID, bookingID)
- booking_flight(bookingID, flightID)

Transaction 3:

- payment(paymentID, amount, paymentStatus)

SQL Commands and Table Specifications

SQL Commands for Selected Transactions/Queries

Transaction 1:

```
SELECT bf.flightID, COUNT(pb.passengerID) AS total_passengers
FROM booking_flight bf
JOIN passenger_booking pb ON bf.bookingID = pb.bookingID
WHERE bf.flightID = 600
GROUP BY bf.flightID;
```

Transaction 2:

```
SELECT 'STA00' + CAST(bf.flightID AS NVARCHAR) AS [Flight ID],
       p.passengerID AS 'Passenger ID', p.fullName AS 'Name', p.email
FROM passenger p
JOIN passenger_booking pb ON p.passengerID = pb.passengerID
JOIN booking_flight bf ON pb.bookingID = bf.bookingID
WHERE bf.flightID = 600 OR bf.flightID = 601
ORDER BY bf.flightID;
```

Transaction 3:

```
SELECT SUM(p.amount) AS total_revenue
FROM payment p
WHERE p.paymentStatus = 'PAID';
```

Table Specifications (Number of Records, Record Sizes)

Transaction 1:

Table	No. of Records	Est. Record Size	Primary Key(s)
passenger	800	~120 bytes	passengerID
booking	60	~4 bytes	bookingID
flight	251	~220 bytes	flightID
passenger_booking	155	~8 bytes	(passengerID, bookingID)
booking_flight	60	~8 bytes	(bookingID, flightID)

Transaction 2:

Table	No. of Records	Est. Record Size	Primary Key(s)
passenger	800	~120 bytes	passengerID
booking	60	~4 bytes	bookingID
flight	251	~220 bytes	flightID
passenger_booking	155	~8 bytes	(passengerID, bookingID)
booking_flight	60	~8 bytes	(bookingID, flightID)

Transaction 3:

Table	No. of Records	Est. Record Size	Primary Key(s)
payment	60	~8 bytes	paymentID

Pre-Improvement Transaction Analysis

Performance Results from RDBMS Query Analyzer

Transaction 1: Count Total Passengers on a specific flight

```


28
29  -- Count Total Passengers on Specific Flight
30  set statistics time on
31  < SELECT
32    ...bf.flightID, COUNT(DISTINCT pb.passengerID) AS total_passengers
33  FROM dbo.booking_flight bf
34  JOIN dbo.passenger_booking pb
35  ON bf.bookingID = pb.bookingID
36  WHERE bf.flightID = 600
37  GROUP BY bf.flightID;
38  set statistics time off
39
40
41  -- List all passengers on specific flight

```

Results Messages

8:28:35 PM Started executing query at Line 29
(1 row affected)

SQL Server Execution Times:
CPU time = 4 ms, elapsed time = 3 ms.
Total execution time: 00:00:00.037



Properties	
SELECT	
Name	Value
Statement	SELECT bf.flightID, COUNT(DISTINCT p...
Cached plan size	40 KB
Estimated Number of Rows Per Execution	1
Estimated Number of Rows for All Executions	0
> Set Options	ANSI_NULLS: True, ANSI_PADDING: True, ...
Optimization Level	FULL
Reason For Early Termination Of Statement Opti...	Good Enough Plan Found
CardinalityEstimationModelVersion	150
QueryHash	0xF32EC966FF0FACB3
QueryPlanHash	0x535D5C895B15D49A
RetrievedFromCache	true
SecurityPolicyApplied	False
> MemoryGrantInfo	
> OptimizerHardwareDependentProperties	
> OptimizerStatsUsage	
> TraceFlags	
CompileTime	21
CompileCPU	21
CompileMemory	376

These pictures represent the result from the current version in which index structure has not been implemented yet. As can be seen in the Query Plan, It still uses the clustered index scan. Execution time: CPU Time = 4ms and elapsed time = 3 ms; Result in Total execution time: 00:00:00:037

Transaction Analysis Form for Transaction 1		
April 28, 2025		
Transaction	Count Total Passengers on Specific Flight	
Volume	Average	Depends on how many flight at that specific moment in time But best guess is around 50 per hour
	Peak	May be hundreds or more
<pre> SELECT bf.flightID, COUNT(DISTINCT pb.passengerID) AS total_passengers FROM dbo.booking_flight bf JOIN dbo.passenger_booking pb ON bf.bookingID = pb.bookingID WHERE bf.flightID = 600 GROUP BY bf.flightID; </pre>		<p>Predicate</p> <pre>bf.flightID = 600</pre> <p>Join Attribute</p> <pre> FROM dbo.booking_flight bf JOIN dbo.passenger_booking pb ON bf.bookingID = pb.bookingID </pre> <p>Ordering Attribute</p> <p>-</p> <p>Grouping Attribute</p> <pre>GROUP BY bf.flightID;</pre> <p>Built-in functions</p> <p>-</p> <p>Attributes Updated</p> <p>-</p>


```

graph LR
    booking_flight[booking_flight] -- 1 --> passenger_booking[passenger_booking]
    booking_flight -- "is for" --> passenger_booking
    
```

The diagram illustrates the relationship between two tables: **booking_flight** and **passenger_booking**. The **booking_flight** table has a primary key **bookingID** (int, not null). The **passenger_booking** table has a primary key **passengerID** (int) and a foreign key **bookingID** (int, not null). A relationship line connects the **bookingID** in **booking_flight** to the **bookingID** in **passenger_booking**, labeled "is for". The multiplicity "1" is associated with the **bookingID** in **booking_flight**, and "0..*" is associated with the **bookingID** in **passenger_booking**.

Access	Entity	Type of	No. of References		
			Per Transactions	Per Transaction	Peak Per Hour*
1	Booking_flight	Read	4	4	400 or more
2	passenger_booking	Read	155	155	155000 or more
Total References			159	159	155400 or more

The Number of references taken from the Query Plan number.

- For the Peak per hour the number per transaction multiplied with the volume.
One big note is that it really depends on how big the tables are.

Transaction 2: List all passengers on specific flight

```


40
41 -- List all passengers on specific flight
42
43 set statistics time on
44
45 SELECT 'STA00' + CAST(bf.flightID as NVARCHAR) as [Flight ID],
46 ... p.passengerID as 'Passenger ID', p.fullName as 'Name', p.email
47 FROM dbo.passenger..... AS p
48 JOIN dbo.passenger_booking AS pb ON p.passengerID = pb.passengerID
49 JOIN dbo.booking_flight AS bf ON pb.bookingID = bf.bookingID
50 WHERE bf.flightID = 606 or bf.flightID = 601
51 ORDER BY bf.flightID;
52
53 set statistics time off
54
55
56 -- Calculate total revenue from complete payments
57 set statistics time on

```

Results Messages

8:29:04 PM Started executing query at Line 43
(7 rows affected)

SQL Server Execution Times:
CPU time = 4 ms, elapsed time = 4 ms.
Total execution time: 00:00:00.012



Properties

Name	Value
Statement	SELECT 'STA00' + CAST(bf.flightID as N...
Cached plan size	40 KB
Estimated Number of Rows Per Execution	5.16667
Estimated Number of Rows for All Execu...	0
> Warnings	Type conversion in expression (CONVER...
> Set Options	ANSI_NULLS: True, ANSI_PADDING: Tru...
Optimization Level	FULL
Reason For Early Termination Of Statem...	Good Enough Plan Found
CardinalityEstimationModelVersion	150
QueryHash	0xDA35EEC3A3E9FDA2
QueryPlanHash	0x0D29C2FB3E03C621
RetrievedFromCache	false
SecurityPolicyApplied	False
> MemoryGrantInfo	
> OptimizerHardwareDependentProperties	
> OptimizerStatsUsage	
> TraceFlags	
CompileTime	17
CompileCPU	17
CompileMemory	496

These pictures represent the result from the current version in which index structure has not been implemented yet. As can be seen in the Query Plan, It still uses the clustered index scan. As can be seen in the Query Plan, It still uses the clustered index scan. Execution time: CPU Time = 4ms and elapsed time = 4 ms; Result in Total execution time: 00:00:00:012

Transaction Analysis Form for Transaction 2			
April 28, 2025			
Transaction	List all passengers on specific flight		
Volume	Average	~ 100 queries at specific point in time	
	Peak	~ 200 or more at specific point in time	
<pre> SELECT 'STA00' + CAST(bf.flightID as NVARCHAR) as [Flight ID], p.passengerID as 'Passenger ID', p.fullName as 'Name', p.email FROM dbo.passenger AS p JOIN dbo.passenger_booking AS pb ON p.passengerID = pb.passengerID JOIN dbo.booking_flight AS bf ON pb.bookingID = bf.bookingID WHERE bf.flightID = 606 or bf.flightID = 601 ORDER BY bf.flightID; </pre>	Predicate	bf.flightID = 606 or bf.flightID = 601	
	Join Attribute	$ \begin{aligned} &\text{FROM dbo.passenger AS } p \\ &\text{JOIN dbo.passenger_booking AS } pb \text{ ON} \\ &p.passengerID = pb.passengerID \\ &\text{JOIN dbo.booking_flight} \\ &\text{AS bf ON pb.bookingID = bf.bookingID} \end{aligned} $	
	Ordering Attribute	bf.flightID	
	Grouping Attribute	-	
	Built-in functions	CAST(bf.flightID as NVARCHAR) as [Flight ID]	
	Attributes Updated	-	


```

classDiagram
    class booking_flight {
        bookingID int PK
        flightID int not null
    }
    class passenger_booking {
        passengerID int PK
        bookingID int PK
    }
    class passenger {
        passengerID int PK
        fullName nchar(200) not null
        phoneNum nchar(20)
        email nchar(100) unique
    }
    booking_flight "0..1" *-- "0..*"
    booking_flight "*" -- "1..1" passenger_booking : is for
    booking_flight "0..*"
    passenger_booking "1..1" -- "1..1" passenger : book

```

Entity	Type of	No. of References

Access			Per Transactions	Per Transaction	Peak Per Hour*
1	passenger	Read	2	2	400 or more
2	passenger_booking	Read	310	310	62000 or more
3	Booking_flight	Read	7	7	1400 or more
Total References			319	319	63800 or more

The Number of references taken from the Query Plan number.

For the Peak per hour, the number per transaction multiply with the volume.

One big note is that it really depends on how big tables are.

Transaction 3: Calculate the total revenue from complete payments

```


54
55
56 -- Calculate total revenue from complete payments
57 set statistics time on
58
59 select sum(p.amount) as total_revenue
60 from payment p
61 where p.paymentStatus = 'PAID';
62
63 set statistics time off
64

```


Results Messages

8:30:07 PM Started executing query at Line 56
SQL Server parse and compile time:
CPU time = 0 ms, elapsed time = 0 ms.
(1 row affected)

SQL Server Execution Times:
CPU time = 2 ms, elapsed time = 2 ms.
Total execution time: 00:00:00.025



These pictures represent the result from the current version, in which the index structure has not yet been implemented. As seen in the Query Plan, the clustered index scan is still used. As seen in the Query Plan, It still uses the clustered index scan. SQL Server Execution time: CPU Time = 2ms and elapsed time = 2 ms; Result in Total execution time: 00:00:00:025

Transaction Analysis Form for Transaction 3					
April 28, 2025					
Transaction	Calculate total revenue from completed payment				
Volume	Average	5 from time to time because this SQL script is not commonly used a lot.			
	Peak	At most 10			
<pre>select sum(p.amount) as total_revenue from payment p where p.paymentStatus = 'PAID';</pre>		Predicate	p.paymentStatus = 'PAID'		
		Join Attribute	-		
		Ordering Attribute	-		
		Grouping Attribute	-		
		Built-in functions	sum(p.amount)		
		Attributes Updated	-		
 <p>It's just 1 table.</p>					
Access	Entity	Type of	No. of References		
			Per Transactions	Per Transaction	Peak Per Hour
1	payment	Read	23	23	23*10= 230
Total References			23	23	230

Post-Improvement Transaction Analysis

```

-- update
-- Speed up filtering booking_flight by flightID
CREATE NONCLUSTERED INDEX IX_booking_flight_flightID
    ON dbo.booking_flight(flightID)
    INCLUDE (bookingID);

-- Speed up joining passenger_booking on bookingID
CREATE NONCLUSTERED INDEX IX_passenger_booking_bookingID
    ON dbo.passenger_booking(bookingID)
    INCLUDE (passengerID);

-- Speed up summing payments for PAID only-filtered index
CREATE NONCLUSTERED INDEX IX_payment_paid_amount
    ON dbo.payment(amount)
    WHERE paymentStatus = 'PAID';

-- Speed up the passenger
CREATE NONCLUSTERED INDEX IX_passenger_Cover
    ON dbo.passenger(passengerID)
    INCLUDE (fullName, email);

```

We Improved by transforming the current clustered index (ordered) of each transaction as have been seen in the query plan in the pre improvement section into nonclustered index which would allow the transaction to be speed up with the help of pointer from the non clustered index, we also add the nonclustered index into the passenger to faster passenger table up. Additionally, we updated the sql to allow even faster response time.

```

SELECT
    bf.flightID,
    COUNT_BIG(*) AS total_passengers -- COUNT_BIG for huge tables
FROM dbo.booking_flight AS bf
JOIN dbo.passenger_booking AS pb
    ON bf.bookingID = pb.bookingID
WHERE bf.flightID = 600
GROUP BY bf.flightID;

```

To handle if transaction getting larger by using COUNT_BIG(*)

```

SELECT
    FlightID      = 'STA00' + CONVERT(NVARCHAR(10), bf.flightID),
    pb.passengerID,
    p.fullName,
    p.email
FROM dbo.booking_flight AS bf
JOIN dbo.passenger_booking AS pb
    ON bf.bookingID = pb.bookingID
JOIN dbo.passenger AS p
    ON pb.passengerID = p.passengerID
WHERE bf.flightID IN (601, 606)      -- cleaner than OR
ORDER BY bf.flightID;

```

Change the OR to set in (...) because it is cleaner.

```
SELECT
|   total_revenue = SUM(p.amount)
FROM dbo.payment AS p
WHERE p.paymentStatus = 'PAID';      -- uses the filtered index
```

Use the newly implemented index.

Performance Results after Optimization

Transaction 1: Count Total Passengers on a specific flight

```
-- updated sql
SET STATISTICS time ON;


SELECT
    bf.flightID,
    COUNT_BIG(*) AS total_passengers -- COUNT_BIG for huge tables
FROM dbo.booking_flight AS bf
JOIN dbo.passenger_booking AS pb
ON bf.bookingID = pb.bookingID
WHERE bf.flightID = 600
GROUP BY bf.flightID;

SET STATISTICS time OFF;
```

Messages

2:49 AM Started executing query at Line 90
(1 row affected)

SQL Server Execution Times:
CPU time = 0 ms, elapsed time = 0 ms.
Total execution time: 00:00:00.005



Properties

SELECT

Name	Value
Statement	SELECT bf.flightID, COUNT_BIG(*) ...
Cached plan size	32 KB
Estimated Number of Rows Per Execution	1
Estimated Number of Rows for All Execut...	0
> Set Options	ANSI_NULLS: True, ANSI_PADDING: Tru...
Optimization Level	FULL
Reason For Early Termination Of Statem...	Good Enough Plan Found
CardinalityEstimationModelVersion	150
QueryHash	0x1262C476C11CF34B
QueryPlanHash	0xC5A850495A5262D8
RetrievedFromCache	false
SecurityPolicyApplied	False
> MemoryGrantInfo	
> OptimizerHardwareDependentProperties	
> OptimizerStatsUsage	
> TraceFlags	
CompileTime	11
CompileCPU	11
CompileMemory	264

From the picture above, the query plan is noticeably shorter than the original one.

Transaction Analysis Form for Transaction 1					
April 28, 2025					
Transaction	Count Total Passengers on Specific Flight				
Volume	Average	Depends on how many flight at that specific moment in time But best guess is around 50 per hour			
	Peak	May be hundreds or more			
<pre> SELECT bf.flightID, COUNT_BIG(*) AS total_passengers -- COUNT_BIG for huge tables FROM dbo.booking_flight AS bf JOIN dbo.passenger_booking AS pb ON bf.bookingID = pb.bookingID WHERE bf.flightID = 600 GROUP BY bf.flightID; </pre>		Predicate	bf.flightID = 600		
		Join Attribute	FROM dbo.booking_flight AS bf JOIN dbo.passenger_booking AS pb ON bf.bookingID = pb.bookingID		
		Ordering Attribute	-		
		Grouping Attribute	GROUP BY bf.flightID;		
		Built-in functions	-		
		Attributes Updated	-		
Access	Entity	Type of	No. of References		
			Per Transactions	Per Transaction	Peak Per Hour
1	Booking_flight	Read	1	1	1*100 = 100 or more
2	passenger_booking	Read	4	4	4*100 = 400 or more

Total References	5	5	500 or more
------------------	---	---	-------------

The Number of references taken from the Query Plan number.

For the Peak per hour the number per transaction multiply with the volume.

One big note is that it really depends on how big tables are.

Transaction 2: List all passengers on specific flight

```


106  -- transaction 2
107  SET STATISTICS time ON;
108
109  SELECT
110    FlightID    = 'STA00' + CONVERT(NVARCHAR(10), bf.flightID),
111    pb.passengerID,
112    p.fullName,
113    p.email
114   FROM dbo.booking_flight AS bf
115  JOIN dbo.passenger_booking AS pb
116  |  ON bf.bookingID = pb.bookingID
117  JOIN dbo.passenger AS p
118  |  ON pb.passengerID = p.passengerID
119  WHERE bf.flightID IN (601, 606)      -- cleaner than OR
120  ORDER BY bf.flightID;
121
122  SET STATISTICS time OFF;
123

```

Results **Messages**

8:54:26 AM Started executing query at Line 106
(7 rows affected)

SQL Server Execution Times:
CPU time = 0 ms, elapsed time = 0 ms.
Total execution time: 00:00:00.009



Properties


SELECT

Filter for any field...

Name	Value
Statement	SELECT FlightID = 'STA00' + CONV...
Cached plan size	40 KB
Degree of Parallelism	1
Estimated Number of Rows Per Execution	5.16667
Estimated Number of Rows for All Execut...	0
> Warnings	Type conversion in expression (CONVER...
> Set Options	ANSI_NULLS: True, ANSI_PADDING: Tru...
Optimization Level	FULL
Reason For Early Termination Of Statem...	Good Enough Plan Found
CardinalityEstimationModelVersion	150
QueryHash	0x0E5006A1B3687BFC
QueryPlanHash	0xFFE479DD36AE9B65
RetrievedFromCache	true
SecurityPolicyApplied	False
> MemoryGrantInfo	
> OptimizerHardwareDependentProperties	
> OptimizerStatsUsage	
> TraceFlags	
> QueryTimeStats	
CompileTime	6
CompileCPU	6
CompileMemory	488

From the picture above, the query plan is slightly shorter than the pre improved one. This version cuts the sort out and transforms the clustered index into the non-clustered index which allows the execution time to be faster than the pre-improve one.

Transaction Analysis Form for Transaction 2		
April 28, 2025		
Transaction	List all passengers on specific flight	
Volume	Average	~ 100 queries at specific point in time
	Peak	~ 200 or more at specific point in time
<pre> SELECT FlightID = 'STA00' + CONVERT(NVARCHAR(10), bf.flightID), pb.passengerID, p.fullName, p.email FROM dbo.booking_flight AS bf JOIN dbo.passenger_booking AS pb ON bf.bookingID = pb.bookingID JOIN dbo.passenger AS p ON pb.passengerID = p.passengerID WHERE bf.flightID IN (601, 606) -- cleaner than OR ORDER BY bf.flightID; </pre>	Predicate <code>bf.flightID IN (601, 606)</code>	
	Join Attribute <code>FROM dbo.booking_flight AS bf JOIN dbo.passenger_booking AS pb ON bf.bookingID = pb.bookingID JOIN dbo.passenger AS p ON pb.passengerID = p.passengerID</code>	
	Ordering Attribute <code>bf.flightID</code>	
	Grouping Attribute <code>-</code>	
	Built-in functions <code>FlightID = 'STA00' + CONVERT(NVARCHAR(10),</code>	
	Attributes Updated <code>-</code>	



Access	Entity	Type of	No. of References		
			Per Transactions	Per Transaction	Peak Per Hour
1	passenger	Read	2	2	400 or more

2	passenger _booking	Read	7	7	1400 or more
3	Booking _flight	Read	7	7	1400 or more
Total References			16	16	3200 or more

The Number of references taken from the Query Plan number.

For the Peak per hour the number per transaction multiply with the volume.

One big note is that it really depends on how big tables are.

Transaction 3: Calculate the total revenue from complete payments

```


125
126  -- transaction 3
127  SET STATISTICS time ON;
128
129  <-- SELECT
130  |   total_revenue = SUM(p.amount)
131  FROM dbo.payment AS p
132  WHERE p.paymentStatus = 'PAID';      -- uses the filtered index
133
134  SET STATISTICS time OFF;
135

```

Results Messages

8:55:42 AM Started executing query at Line 126
SQL Server parse and compile time:
CPU time = 0 ms, elapsed time = 0 ms.
(1 row affected)

SQL Server Execution Times:
CPU time = 1 ms, elapsed time = 0 ms.
Total execution time: 00:00:00.014




Properties

SELECT

Name	Value
Statement	SELECT SUM([p].[amount]) [total_reven...
Cached plan size	24 KB
Degree of Parallelism	1
> Parameter List	@1
Estimated Number of Rows Per Execution	1
Estimated Number of Rows for All Execut...	0
> Warnings	
> Set Options	ANSI_NULLS: True, ANSI_PADDING: Tru...
Optimization Level	FULL
Reason For Early Termination Of Statem...	Good Enough Plan Found
CardinalityEstimationModelVersion	150
QueryHash	0xBE8318CE68655DA5
QueryPlanHash	0x765BD356E569C4C0
RetrievedFromCache	false
SecurityPolicyApplied	False
> UnmatchedIndexes	
> MemoryGrantInfo	
> OptimizerHardwareDependentProperties	
> OptimizerStatsUsage	
> TraceFlags	
> QueryTimeStats	
CompileTime	16
CompileCPU	16
CompileMemory	280

From the picture above, the query plan is similar to the pre improvement but the difference is the index has been implemented to the filter level, which allows the search to be faster, and reduces the execution time.

```
-- Speed up summing payments for PAID only-filtered index
CREATE NONCLUSTERED INDEX IX_payment_paid_amount
    ON dbo.payment(amount)
    WHERE paymentStatus = 'PAID';
```

Transaction Analysis Form for Transaction 3 After improvement								
April 28, 2025								
Transaction	Calculate total revenue from completed payment							
Volume	Average	5 from time to time because this SQL script is not commonly used a lot.						
	Peak	10 or more						
<pre> SELECT total_revenue = SUM(p.amount) FROM dbo.payment AS p WHERE p.paymentStatus = 'PAID'; -- uses the filtered index </pre>		Predicate	total_revenue = SUM(p.amount), p.paymentStatus = 'PAID';					
		Join Attribute	-					
		Ordering Attribute	-					
		Grouping Attribute	-					
		Built-in functions	sum(p.amount)					
		Attributes Updated	-					
								
<p>It's just 1 table.</p>								
Access	Entity	Type of	No. of References					
			Per Transactions	Per Transaction	Peak Per Hour			
1	payment	Read	23	23	230 or more			
Total References			23	23	230 or more			

Comparison with Pre-Improvement Results

Type / Time	Pre-Improvement			Post-Improvement		
	Server Execution CPU / elapsed	Compile Time/CPU	Total Execution	Server Execution CPU / elapsed	Compile Time/CPU	Total Execution
Transaction 1	4ms/ 3ms	21ms/21ms	00:00:00.037	0ms/0ms	11ms/11ms	00:00:00.005
Transaction 2	4ms/4ms	17ms/17ms	00:00:00.012	0ms/0ms	5ms/5ms	00:00:00.009
Transaction 3	2ms/2ms	6ms/6ms	00:00:00.025	1ms/0ms	1ms/1ms	00:00:00.014

Conclusion

After implementing indexes, it does seem to be faster across all transactions and make look up and search from related tables faster. One big limitation is Data size. We have limited data size so the improvement that we made might not seem to be a big difference if we compare both pre-improvement and post-improvement together.

Discussion

The analysis of all three selected transactions counting passengers on a specific flight, listing passengers, and calculating total revenue from completed payments shows a marked improvement in performance after applying database tuning techniques. Initially, all transactions relied on clustered index scans, leading to suboptimal execution times. After optimization, we implemented non-clustered indexes on frequently filtered and joined columns (e.g., flightID, passengerID, paymentStatus). This significantly reduced the CPU and elapsed execution times.

Transaction 1 (Passenger Count) benefited from using COUNT_BIG(*) and optimized join paths, reducing execution time from 37ms to 5ms.

Transaction 2 (Passenger List) saw improvements by replacing OR with IN, avoiding unnecessary sorts, and leveraging the new indexes, reducing execution time from 12ms to 9ms.

Transaction 3 (Revenue Calculation) gained from a filtered index on paymentStatus, bringing the execution time down from 25ms to 14ms.

The improvement is consistent across all performance metrics: CPU time, elapsed time, and total execution time. This validates the effectiveness of the applied tuning strategies.

However, it is important to note that real-world performance can vary with larger data volumes. Future improvements may include partitioning, query refactoring, and materialized views for extremely high-traffic systems.

References

MikeRayMSFT, “Clustered and nonclustered indexes - SQL Server,” *Microsoft Learn*.
<https://learn.microsoft.com/en-us/sql/relational-databases/indexes/clustered-and-nonclustered-indexes-described?view=sql-server-ver16>